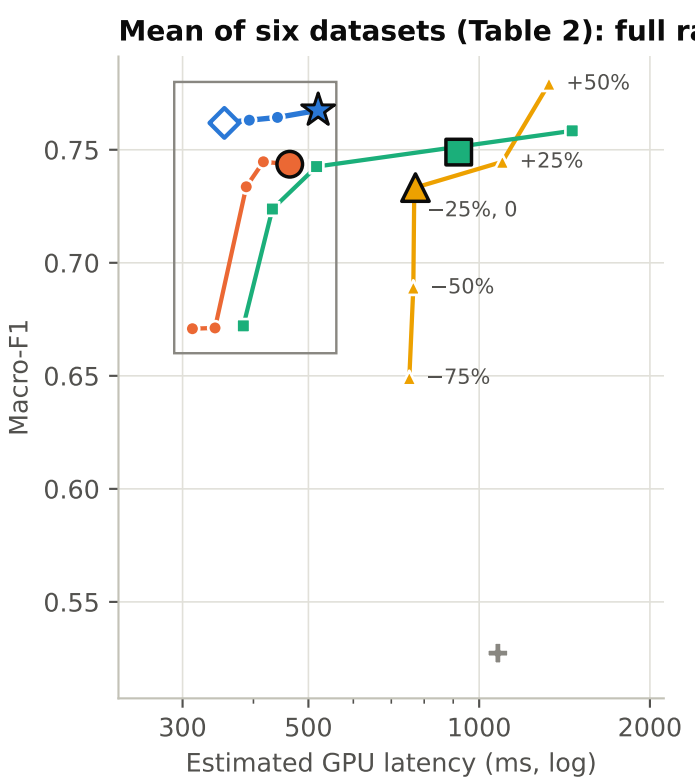

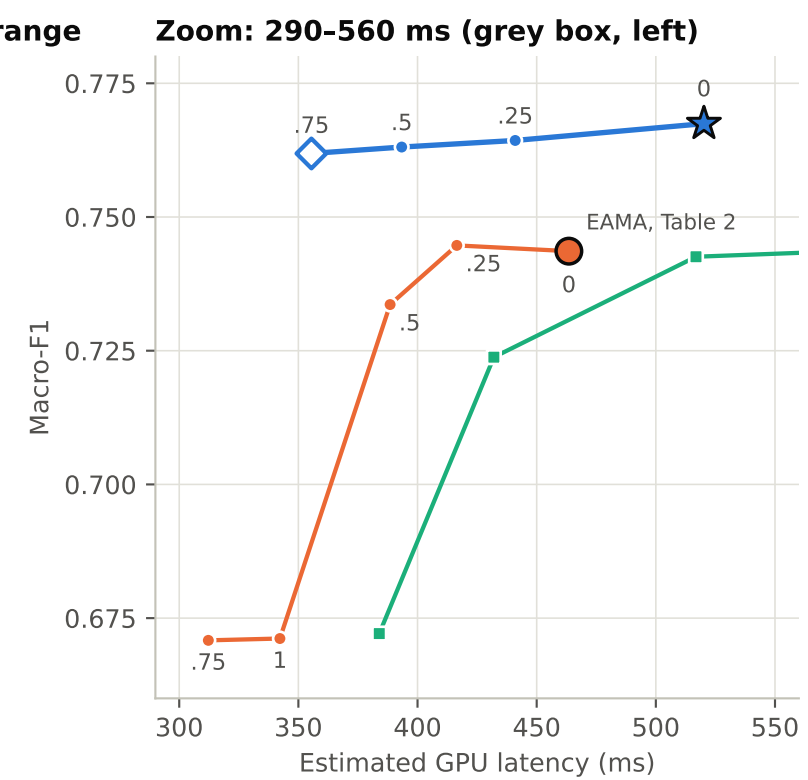

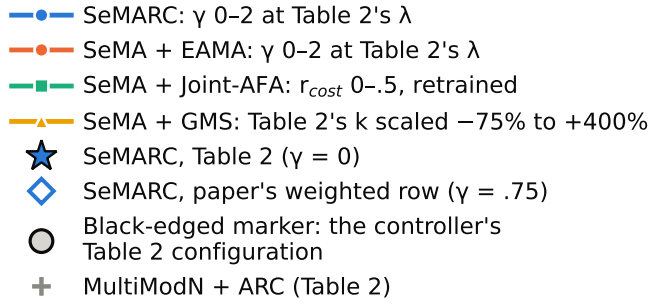

Mean of the six datasets. Each knob value is shared by all six; lines are each method's Pareto front over its knob. SeMARC and EAMA: γ varied at Table 2's λ (EAMA: each run's validation-picked λ); EAMA's modality costs scale as SeMARC's do, $z_j = t_j^{\gamma} / \Sigma\, t^{\gamma}$. GMS: every dataset's Table 2 k scaled by the same percentage (rounded, $1 \le k \le M$). EAMA's and GMS's lines also run through their Table 2 points (EAMA's γ = 0 is beaten by its γ = .25). Numbers next to the dots give γ or the change in k. Joint-AFA's black-edged marker is the paper's run.