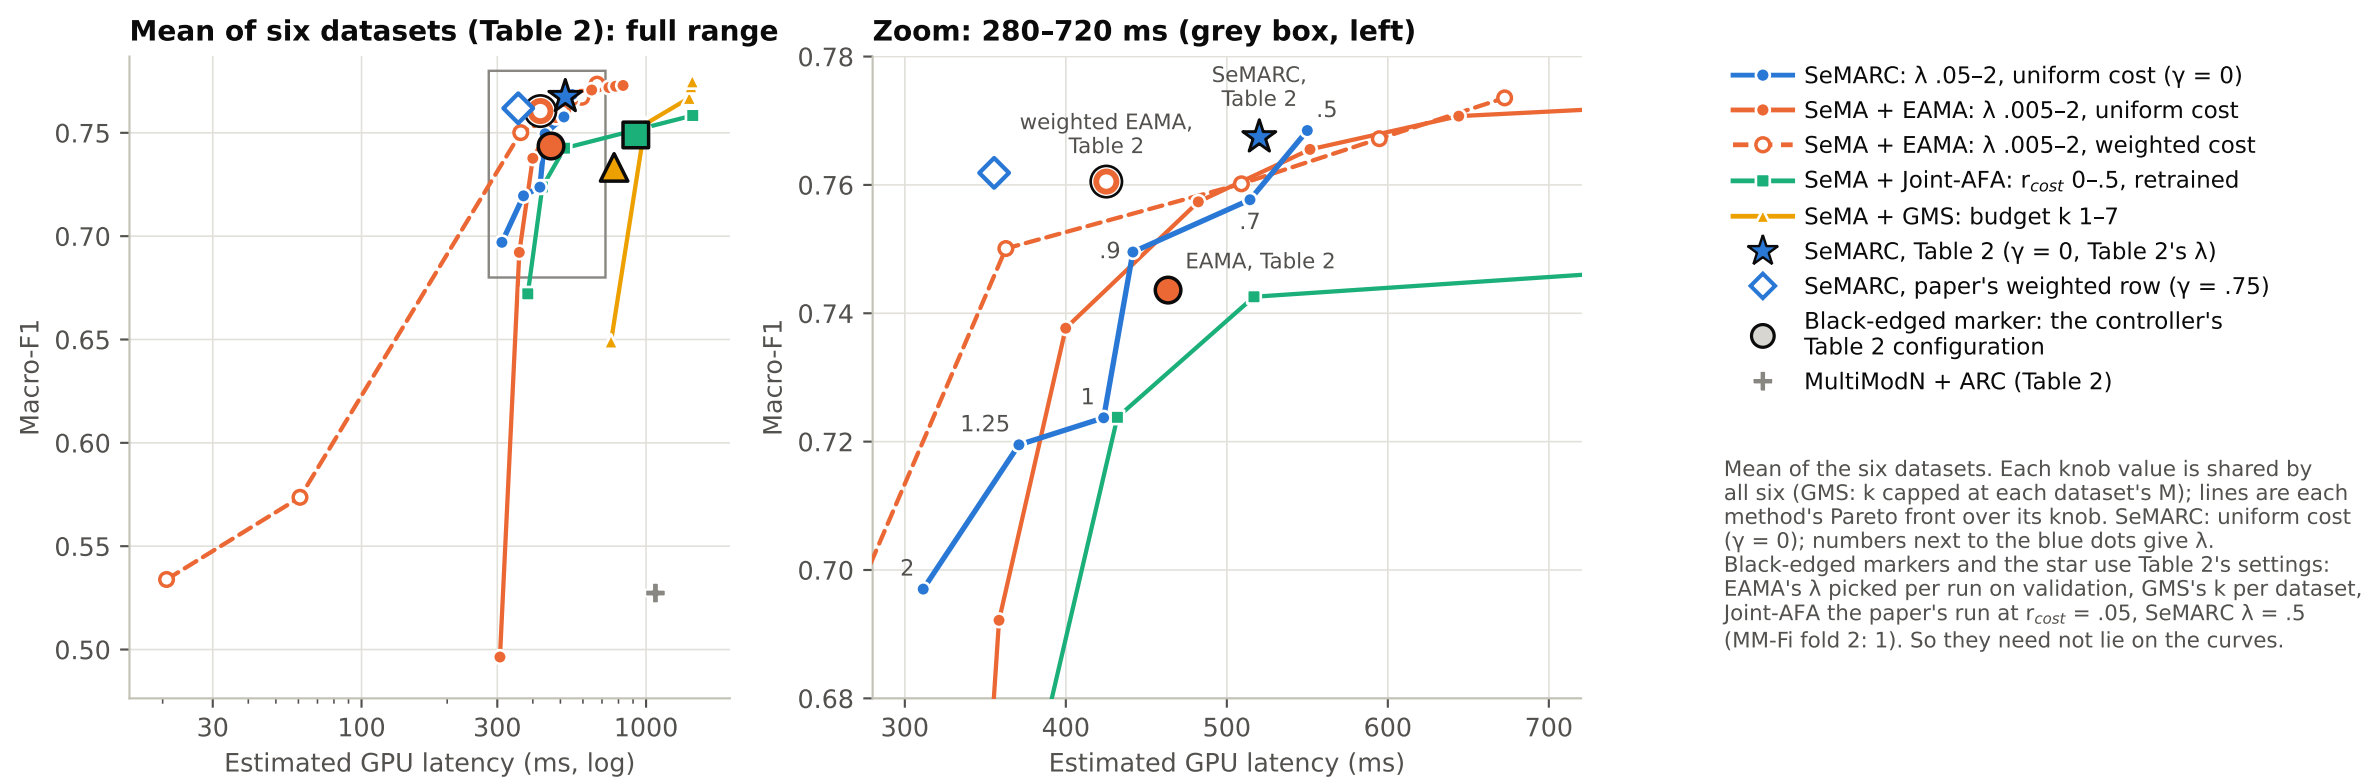

Mean of the six datasets. Each knob value is shared by all six (GMS: k capped at each dataset's M); lines are each method's Pareto front over its knob. SeMARC: uniform cost (γ = 0); numbers next to the blue dots give λ.
Black-edged markers and the star use Table 2's settings: EAMA's λ picked per run on validation, GMS's k per dataset, Joint-AFA the paper's run at r_{cost} = .05, SeMARC λ = .5 (MM-Fi fold 2: 1). So they need not lie on the curves.